

## SQL vs. NoSQL: Key Differences and Use Cases

SQL databases are **ideal for structured data** with a well-defined **schema**. They follow the principles of **ACID transactions** (Atomicity, Consistency, Isolation, Durability), making them suitable for applications requiring strong consistency, such as banking systems and enterprise software.

On the other hand, **NoSQL (Not Only SQL)** databases provide a **flexible schema** and come in various types, including **document stores, key-value stores, column-family stores, and graph databases**. NoSQL databases are particularly advantageous when **minimal latency** and **high scalability** are required.

## Why NoSQL for Low-Latency Operations?

NoSQL databases are optimized for fast read and write operations. Consider **Redis**, a popular in-memory key-value store:

- It uses **hashing mechanisms** that allow  **$O(1)$  or  $O(\log n)$  time complexity** for lookups, making it significantly faster than traditional SQL queries that may involve multiple joins.
- Since NoSQL databases often store **denormalized** data, they avoid the complexity of expensive **JOIN operations**, allowing for **faster queries** at the cost of some redundancy.

## Developer-Friendly Approach

- NoSQL databases are **schema-less**, meaning they **do not require migrations** when the data structure changes.
- Unlike SQL databases that rely on ORMs (Object-Relational Mappers) to map relational data to objects, NoSQL databases can store **JSON documents directly**, making them **more natural** for developers working with modern web applications.

## Scalability: SQL vs. NoSQL

- SQL databases typically **scale vertically** (by adding more resources to a single powerful server) and use **read replicas** to distribute read loads. However, **writes are limited to the primary leader**, making write scalability a challenge.
- NoSQL databases are **designed for horizontal scaling**, supporting **sharding** (partitioning data across multiple servers). Each **shard** can reside on a different server, allowing for **distributed writes and reads**, leading to **high availability and fault tolerance**.

Factor	SQL	NoSQL
Schema	Fixed, structured	Flexible, dynamic
Data Relationships	Strong relationships, normalized	Weak relationships, denormalized
Read Performance	Optimized for complex queries	Optimized for fast lookups
Write Performance	Limited by a single leader	Distributed writes with sharding
Scalability	Vertical scaling, read replicas	Horizontal scaling, sharding
Use Cases	Financial systems, ERP, CRM	Real-time analytics, caching, IoT, social media